

Final Exam - GuardRail and Shadow Test

Sajib (RYHAN) Suny
N01654285

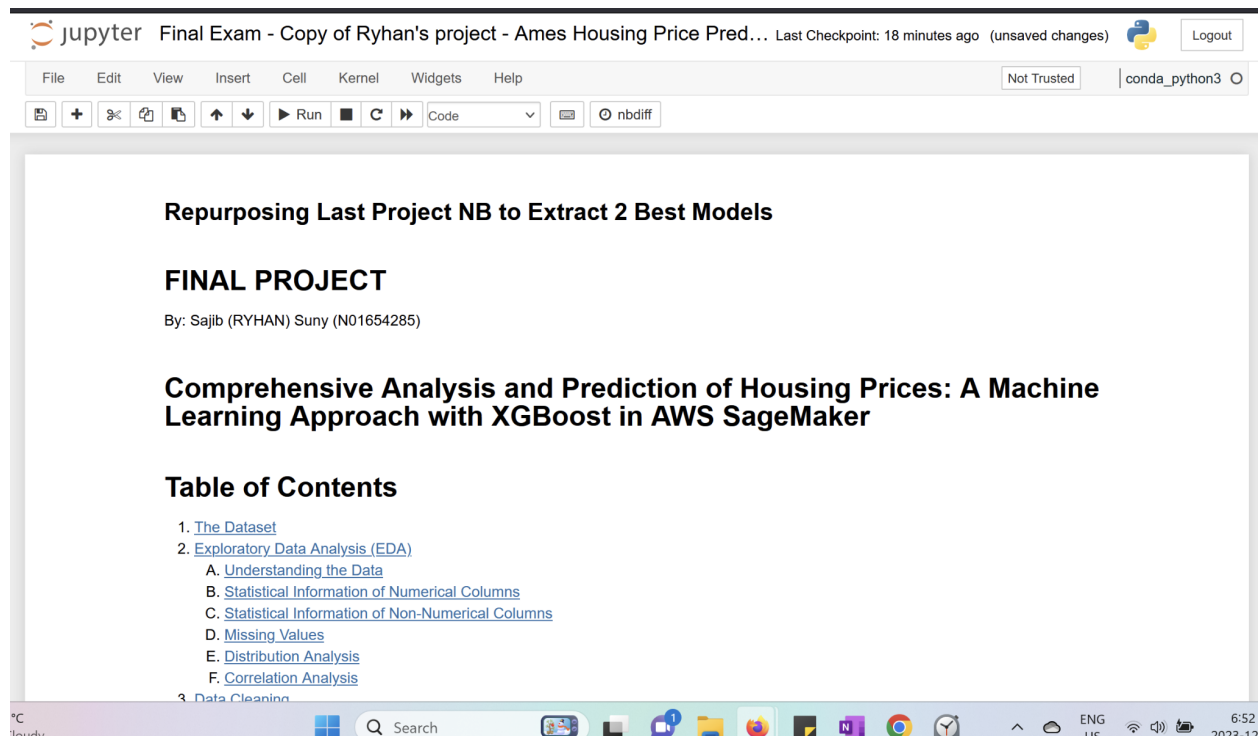
**** Prompt: You will simulate a deployment to production through SageMaker guardrail and SageMaker Shadow testing**

You will start from a data set and prepare the data for your project algorithm (2marks). You should show me the data set and algorithm you have used

1. GuardRail

Showing DATASET and Algorithm

Source: Previous Project Notebook



- Below Screenshot shows dataset - Ames housing prices

Reading the Dataset

```
In [1]: import pandas as pd
import numpy as np

dataset = pd.read_csv("AmesHousing.csv")

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/pandas/core/computation/expressions.py:21: UserWarning: Pand
as requires version '2.8.0' or newer of 'numexpr' (version '2.7.3' currently installed).
from pandas.core.computation.check import NUMEXPR_INSTALLED
```

```
In [3]: # Display basic information about the dataset
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2930 entries, 0 to 2929
Data columns (total 82 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Order               2930 non-null   int64
1   PID                 2930 non-null   int64
2   MS SubClass         2930 non-null   int64
3   MS Zoning            2930 non-null   object
4   Lot Frontage        2440 non-null   float64
5   Lot Area            2930 non-null   int64
6   Street              2930 non-null   object
7   Alley               198 non-null    object
8   Lot Shape           2930 non-null   object
9   Land Contour        2930 non-null   object
10  Utilities            2930 non-null   object
11  Lot Config          2930 non-null   object
12  Land Slope           2930 non-null   object
13  Neighborhood         2930 non-null   object
```

- Below screenshots show the normal Training Job configuration and job name, please note the last digits of the job name: **18-34** as an identifier.

```

In [91]: XGB_job = "FinalProject-Ryhan-XGBoost-" + time.strftime("%Y-%m-%d-%H-%M-%S", time.gmtime())

print("Job name is:", XGB_job)

xgb_training_params = {
    "TrainingJobName": XGB_job,
    "AlgorithmSpecification": {
        "TrainingImage": "container",
        "TrainingInputMode": "File"
    },
    "RoleArn": role,
    "ResourceConfig": {"InstanceCount": 1, "InstanceType": "ml.c4.2xlarge", "VolumeSizeInGB": 10},
    "HyperParameters": {
        "max_depth": "5",           # Maximum depth of a tree. Increasing this value will make the model more complex and more likely to overfit.
        "eta": "0.2",              # Step size shrinkage used in updates to prevent overfitting. After each boosting step, we can directly get the maximum improvement on the objective, which is then multiplied by eta to form the step size.
        "gamma": "4",              # Minimum loss reduction required to make a further partition on a leaf node of the tree.
        "min_child_weight": "6",    # Minimum sum of instance weight (hessian) needed in a child.
        "verbosity": "0",           # Verbosity of printing messages. Valid values are 0 (silent), 1 (warning), 2 (info), and 3 (debug).
        "objective": "reg:squarederror", # Objective function for regression.
        # Removed num_class as it's not relevant for regression
        "num_round": "10"           # The number of rounds for boosting
    },
    "StoppingCondition": {
        "MaxRuntimeInSeconds": 86400
    },
    "OutputDataConfig": {
        "S3OutputPath": f"s3://{bucket}/{prefix}/xgboost"
    },
    "InputDataConfig": [
        {
            "ChannelName": "train",
            "DataSource": {

```

```

                "S3DataSource": {
                    "S3DataType": "S3Prefix",
                    "S3Uri": f"s3://{bucket}/{prefix}/train/",
                    "S3DataDistributionType": "FullyReplicated"
                },
            "ContentType": "application/x-recordio-protobuf",
            "CompressionType": "None",
            "RecordWrapperType": "None"
        },
        {
            "ChannelName": "validation",
            "DataSource": {
                "S3DataSource": {
                    "S3DataType": "S3Prefix",
                    "S3Uri": f"s3://{bucket}/{prefix}/validation/",
                    "S3DataDistributionType": "FullyReplicated"
                },
            "ContentType": "application/x-recordio-protobuf",
            "CompressionType": "None",
            "RecordWrapperType": "None"
        }
    ]
}

```

Job name is: FinalProject-Ryhan-XGBoost-2023-11-26-19-18-34

- Here is the normal Training Job with Xgboost Algorithm shown, job completed as shown in output and in GUI below

TRAINING

In [92]: %%time

```
sm = boto3.client("sagemaker")

sm.create_training_job(**xgb_training_params)

# Checking training job status
status = sm.describe_training_job(TrainingJobName=XGB_job)["TrainingJobStatus"]
print(status)

sm.get_waiter("training_job_completed_or_stopped").wait(TrainingJobName=XGB_job)
if status == "Failed":
    message = sm.describe_training_job(TrainingJobName=XGB_job)["FailureReason"]
    print("Training failed with the following error: {}".format(message))
    raise Exception("Training job failed")
```

InProgress

CPU times: user 91.1 ms, sys: 5.51 ms, total: 96.6 ms

Wall time: 4min

The screenshot shows the Amazon SageMaker console interface. The left sidebar contains navigation links for Governance, HyperPod Clusters, Ground Truth, Notebook, Processing, Training, and Inference. The main content area displays the details for a specific training job.

Job settings

Job name	Status	SageMaker metrics time series	IAM role ARN
FinalProject-Ryhan-XGBoost-2023-11-26-19-18-34	Completed View history	Disabled	arn:aws:iam::378639026377:role/fast-ai-academic-85-Student-Azure
ARN arn:aws:sagemaker:us-east-1:378639026377:training-job/FinalProject-Ryhan-XGBoost-2023-11-26-19-18-34	Creation time Nov 26, 2023 19:18 UTC	Training time (seconds) 78	
	Last modified time Nov 26, 2023 19:21 UTC	Billable time (seconds) 78	
		Managed spot training savings 0%	
		Tuning job source/parent -	

Amazon SageMaker > Hyperparameter tuning jobs > FinalProj-Ryhan-XGB-26-19-33-53

FinalProj-Ryhan-XGB-26-19-33-53

Stop tuning job

Hyperparameter tuning job summary

Name FinalProj-Ryhan-XGB-26-19-33-53	Status Completed	Approx. total training duration 18 minute(s)
ARN arn:aws:sagemaker:us-east-1:378639026377:hyperparameter-tuning-job/FinalProj-Ryhan-XGB-26-19-33-53	Creation time Nov 26, 2023 19:37 UTC	
	Last modified time Nov 26, 2023 19:47 UTC	

Best training job | Training jobs | Training job definitions | Tuning Job configuration | Tags

Best training job summary

This training job is the best training job for only this hyperparameter tuning job.

Create model

Name FinalProj-Ryhan-XGB-26-19-33-53-005-5a2863a9	Status Completed	Objective metric validation:rmse	Value 20404.921875
--	---------------------	-------------------------------------	-----------------------

Output

S3 model artifact

<s3://final-project-ryhan/sagemaker/Ames-housing-price-prediction/output/FinalProj-Ryhan-XGB-26-19-33-53-005-5a2863a9/output/model.tar.gz>

-> Improved Best Model Artifact = Ending With 3a9

Since this is the best model retrieved through HP Tuning, this model is better than previous normal Training Job

-> Please note: Until this part (collecting models from last project) I used MYAPPS, since last project was on myapps. Moving forward from here I use AWS ACADEMY as instructed.

****NEXT Prompt: Use the notebook I gave as a starting point for simulating a guardrail. You are expected to use parts of that code in your newly created notebook. • The first guardrail is a failed deployment that the guardrail rollbacks automatically (5 marks). You must show me the “E” letters in the notebook.**

GUARDRAIL FAILURE CASE

We invoke the endpoint during the update operation is in progress.

Note : Invoke endpoint in this notebook is in single thread mode, to stop the invoke requests please stop the cell execution

The E's denote the errors generated from the incompatible model version in the canary fleet.

The purpose of the below cell is to simulate errors in the canary fleet. Since the nature of traffic shifting to the canary fleet is probabilistic, you should wait until you start seeing errors. Then, you may proceed to stop the execution of the below cell. If not aborted, cell will run for 600 invocations.

```
[58]: invoke_endpoint(endpoint_name)

Sending test traffic to the endpoint Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23.
Please wait...
.....
.....Exception: An error occurred (ModelError) when calling the InvokeEndpoint
nt operation: Received server error (500) from primary and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
EException: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
EException: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
EException: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
ventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
E..Exception: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
y and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
E.....Exception: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
E..Exception: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
y and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
E..Exception: An error occurred (ModelError) when calling the InvokeEndpoint operation: Received server error (500) from primary
y and could not load the entire response body. See https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logEventViewer:group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
E.....group=aws/sagemaker/Endpoints/Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23 in account 285666138595 for more information.
E.....
```

Wait for the update operation to complete and verify the automatic rollback.

Please note, I added a print statement for each exception (to see what was going on) so that's why instead of just “E” it is showing E then the full exception error.

AS SEEN ABOVE, it starts with Then goes to EEEE, finally rolls back to

Below is the describe_endpoint output showcasing the successful ROLLBACK alarm

```
In [59]: wait_for_endpoint_in_service(endpoint_name)


sm.describe_endpoint(EndpointName=endpoint_name)

Waiting for endpoint in service


Done!

Out[59]: {'EndpointName': 'Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23',
'EndpointArn': 'arn:aws:sagemaker:us-east-1:285666138595:endpoint/ryhans-deployment-guardrails-canary-2023-12-13-00-40-23',
'EndpointConfigName': 'Ryhans-FinalExam-EpConfig-1-2023-12-13-00-37-25',
'ProductionVariants': [{'VariantName': 'AllTraffic',
'DeployedImages': [{'SpecifiedImage': '683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost:0.90-2-cpu-py3',
'ResolvedImage': '683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost@sha256:0d098653fff2915993d61180da0cde0ed982805093463d40f30212b8050486f18',
'ResolutionTime': datetime.datetime(2023, 12, 13, 0, 40, 24, 649000, tzinfo=tzlocal())}],
'CurrentWeight': 1.0,
'DesiredWeight': 1.0,
'CurrentInstanceCount': 3,
'DesiredInstanceCount': 3}],
'EndpointStatus': 'InService',
'FailureReason': 'One or more configured alarm for automatic rollback deployment is in ALARM state: [TestAlarm-ModelLatency-Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23].',
'CreationTime': datetime.datetime(2023, 12, 13, 0, 40, 23, 976000, tzinfo=tzlocal()),
'LastModifiedTime': datetime.datetime(2023, 12, 13, 0, 58, 41, 165000, tzinfo=tzlocal()),
'LastDeploymentConfig': {'BlueGreenUpdatePolicy': {'TrafficRoutingConfiguration': {'Type': 'CANARY',
'WaitIntervalInSeconds': 120,
'CanarySize': {'Type': 'INSTANCE_COUNT', 'Value': 1}},
'TerminationWaitInSeconds': 120,
'MaximumExecutionTimeoutInSeconds': 1800},
'AutoRollbackConfiguration': {'Alarms': [{'AlarmName': 'TestAlarm-5XXErrors-Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23'},
{'AlarmName': 'TestAlarm-ModelLatency-Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23'}]}},
'ResponseMetadata': {'RequestId': '4778639a-c46f-4921-9c8a-39a4666db708',
'HTTPStatusCode': 200,
'HTTPHeaders': {'x-amzn-requestid': '4778639a-c46f-4921-9c8a-39a4666db708',
'content-type': 'application/x-amz-json-1.1',
'content-length': '1468',
'date': 'Wed, 13 Dec 2023 01:01:07 GMT'},
'RetryAttempts': 0}}
```

- Invocation metrics for failure+rollback case
- Model 1 (ep_config-1) vs model 2 (epConfig-2) vs all



- Below we also see 5xx errors and latency average rising due to faulty model then falling after rollback



****The second guardrail is a successful deployment, and you will show that the guardrail allows it to be deployed successfully (5 marks)**

GUARDRAIL SUCCESS CASE

- Below we see output for success case is (no errors) after updating to ep-onfig-3 (best HP tuned model or model 2 with correct version)

invoke the endpoint during the update operation is in progress.

```
In [63]: invoke_endpoint(endpoint_name, max_invocations=500)
```

Sending test traffic to the endpoint Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23.
Please wait...

Wait for the update operation to complete:

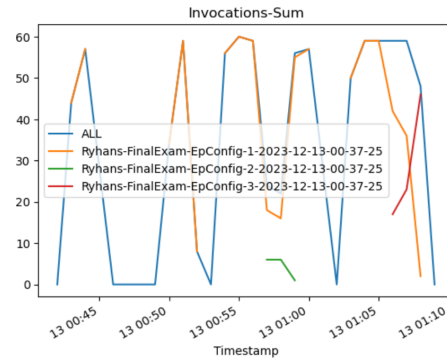

- Output showing successful transition to ep-config-3

```
In [64]: #wait_for_endpoint_in_service(endpoint_name)
sm.describe_endpoint(EndpointName=endpoint_name)

Out[64]: {'EndpointName': 'Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23',
'EndpointArn': 'arn:aws:sagemaker:us-east-1:285666138595:endpoint/ryhans-deployment-guardrails-canary-2023-12-13-00-40-23',
'EndpointConfigName': 'Ryhans-FinalExam-EpConfig-3-2023-12-13-00-37-25',
'ProductionVariants': [{'VariantName': 'AllTraffic',
'DeployedImages': ['683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost:0.90-2-cpu-py3',
'ResolvedImage': '683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost@sha256:0d098653ff2915993d61180da0cde0ed982805093463d40f30212b8050486f18',
'ResolutionTime': datetime.datetime(2023, 12, 13, 1, 3, 2, 137000, tzinfo=tzlocal())],
'CurrentWeight': 1.0,
'DesiredWeight': 1.0,
'CurrentInstanceCount': 3,
'DesiredInstanceCount': 3}],
'EndpointStatus': 'InService',
'CreationTime': datetime.datetime(2023, 12, 13, 0, 40, 23, 976000, tzinfo=tzlocal()),
'LastModifiedTime': datetime.datetime(2023, 12, 13, 1, 9, 53, 370000, tzinfo=tzlocal()),
'LastDeploymentConfig': {'BlueGreenUpdatePolicy': {'TrafficRoutingConfiguration': {'Type': 'CANARY',
'WaitIntervalInSeconds': 120,
'CanarySize': {'Type': 'INSTANCE_COUNT', 'Value': 1}},
'TerminationWaitInSeconds': 120,
'MaximumExecutionTimeoutInSeconds': 1800},
'AutoRollbackConfiguration': {'Alarms': [{'AlarmName': 'TestAlarm-5XXErrors-Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23'},
{'AlarmName': 'TestAlarm-ModelLatency-Ryhans-Deployment-Guardrails-Canary-2023-12-13-00-40-23'}]}],
'ResponseMetadata': {'RequestId': 'b4ed9d3a-21a7-42e5-91a7-e7fbbc66a15a',
'HTTPStatusCode': 200,
'HTTPHeaders': {'x-amzn-requestid': 'b4ed9d3a-21a7-42e5-91a7-e7fbbc66a15a',
'content-type': 'application/x-amz-json-1.1',
'content-length': '1285',
'date': 'Wed, 13 Dec 2023 01:09:59 GMT'},
'RetryAttempts': 0}}
```

Collect the endpoint metrics during the deployment:

- Invocation Metrics



2. Shadow Test

**** I showed you how to use the Shadow test in the console. You can do the same through code. Here is the link:**

https://sagemaker-examples.readthedocs.io/en/latest/sagemaker-shadow-variant/Shadow_variant.html (5 marks)

• You will use the same models you have created in the guardrail. You will design the test in a way that the shadow variant looks better than production variant and you allow that to replace the production variant. (4 marks)

Shadow Test using code

Creating Variants

```
In [40]: from sagemaker.session import production_variant

variant1 = production_variant(
    model_name=model_name,
    instance_type="ml.m5.xlarge",
    initial_instance_count=1,
    variant_name='Variant1',
    initial_weight=1)
variant2 = production_variant(
    model_name=model_name2,
    instance_type="ml.m5.xlarge",
    initial_instance_count=1,
    variant_name='Variant2',
    initial_weight=1)
```

- Invoking endpoint with both variants to compare which model is better, so we can make a decision to promote the shadow or not

endpoint

```
In [42]: endpoint_name = f"Ryhans-ShadowTest-xgb-pred-{datetime.datetime.now():%Y-%m-%d-%H-%M-%S}"
print(f"EndpointName={endpoint_name}")

sm_session.endpoint_from_production_variants(
    name=endpoint_name,
    production_variants=[variant1, variant2])
```

```
EndpointName=Ryhans-ShadowTest-xgb-pred-2023-12-13-04-08-46
----
```

```
Out[42]: 'Ryhans-ShadowTest-xgb-pred-2023-12-13-04-08-46'
```

Invoke the Deployed Models

```
In [44]: import time

# Assuming sm_runtime is already defined
print(f"Sending test traffic to the endpoint {endpoint_name}. \nPlease wait...")

with open('test_data/test_sample_tail_input_cols.csv', 'r') as f:
    for row in f:
        print(".", end="", flush=True)
        payload = row.rstrip('\n')
        sm_runtime.invoke_endpoint(
            EndpointName=endpoint_name,
            ContentType="text/csv",
            Body=payload)
        time.sleep(0.5)

print("Done!")
```

```
Sending test traffic to the endpoint Ryhans-ShadowTest-xgb-pred-2023-12-13-04-08-46.
Please wait...
```

```
.....Done!
```

- Comparing the performance of each variant

```
def evaluate_performance(y_true, y_pred):  
    accuracy = accuracy_score(y_true, y_pred)  
    precision = precision_score(y_true, y_pred, average='binary')  
    recall = recall_score(y_true, y_pred, average='binary')  
    f1 = f1_score(y_true, y_pred, average='binary')  
  
    print(f"Accuracy: {accuracy:.4f}")  
    print(f"Precision: {precision:.4f}")  
    print(f"Recall: {recall:.4f}")  
    print(f"F1 Score: {f1:.4f}")  
  
print("Performance metrics for Variant1:")  
evaluate_performance(labels, y_pred_variant1)  
  
print("\nPerformance metrics for Variant2:")  
evaluate_performance(labels, y_pred_variant2)
```

Performance metrics for Variant1:
Accuracy: 0.4775
Precision: 0.2619
Recall: 0.4681
F1 Score: 0.3359

Performance metrics for Variant2:
Accuracy: 0.4835
Precision: 0.2759
Recall: 0.5106
F1 Score: 0.3582

Variant 2 seems better

so this is the one that we want to use as shadow variant.

- Decision has been made to promote Shadow (model 2)
- Preparing for promotion

Create Endpoint Configuration for Shadow Test

```
51]: create_endpoint_config_response = sm.create_endpoint_config(
    EndpointConfigName="ryhans-shadowtest-endpoint-config",
    ProductionVariants=[
        {
            "VariantName": "Variant1",
            "ModelName": model_name,
            "InstanceType": "ml.m5.xlarge",
            "InitialInstanceCount": 1,
            "InitialVariantWeight": 1
        }
    ],
    ShadowProductionVariants=[
        {
            "VariantName": "Variant2",
            "ModelName": model_name2,
            "InstanceType": "ml.m5.xlarge",
            "InitialInstanceCount": 1, # Added value for InitialInstanceCount
            "InitialVariantWeight": 1
        }
    ]
)
```

```
52]: create_endpoint_response = sm.create_endpoint(
    EndpointName="ryhans-shadowtest-endpoint",
    EndpointConfigName="ryhans-shadowtest-endpoint-config"
)
```

- PROMOTION SUCCESSFUL

Endpoint runtime settings

Update weights

Update instance count

Configure auto scaling

		Variant name ▾	Current weight ▾	Desired weight	Elastic Inference	Instance type ▾	Current instance count
<input type="radio"/>	<div>P</div>	Variant1	0	0	-	ml.m5.xlarge	1
<input type="radio"/>	<div>S</div>	Variant2	100	100	-	ml.m5.xlarge	1

Variants

Identifies a model that you want to host and the resources chosen to deploy for hosting it.

P Production						
Model name	Training job	Variant name	Instance type	Elastic Inference	Initial instance count	Initial weight
Ryhans-SHADOWTEST-xgb-PROD-pred1-2023-12-13-04-06-22	-	Variant1	ml.m5.xlarge	-	1	1

S Shadow						
Model name	Training job	Variant name	Instance type	Elastic Inference	Initial instance count	Initial weight
Ryhans-SHADOWTEST-xgb-UPDATED-pred2-2023-12-13-04-06-22	-	Variant2	ml.m5.xlarge	-	1	1

Showing that Shadow model is performing better

Is shadow variant better?

```
In [61]: import boto3
import pandas as pd
from datetime import datetime, timedelta

# Initialize the CloudWatch client
cloudwatch = boto3.client('cloudwatch')

def get_model_metrics(endpoint_name, variant_name, metric_name, start_time, end_time):
    """ Fetch metrics for a specific model variant from CloudWatch """
    response = cloudwatch.get_metric_data(
        MetricDataQueries=[
            {
                'Id': 'm1',
                'MetricStat': {
                    'Metric': {
                        'Namespace': 'AWS/SageMaker',
                        'MetricName': metric_name,
                        'Dimensions': [
                            {'Name': 'EndpointName', 'Value': endpoint_name},
                            {'Name': 'VariantName', 'Value': variant_name}
                        ]
                    },
                    'Period': 300, # Period in seconds
                    'Stat': 'Average',
                },
                'ReturnData': True,
            },
        ],
        StartTime=start_time,
        EndTime=end_time
    )
    return response['MetricDataResults'][0]['Values']

# Define the metrics to be fetched
metrics = ['Invocation4XXErrors', 'ModelLatency', 'Invocations']

# Set time range for metrics
end_time = datetime.utcnow()
start_time = end_time - timedelta(days=1)
```

```

# Replace with your actual endpoint and variant names
endpoint_name = 'Ryhans-ShadowTest-xgb-pred'
variant1_name = 'Variant1-production'
variant2_name = 'Variant2-shadow'

# Fetch and compare metrics
for metric in metrics:
    variant1_metrics = get_model_metrics(endpoint_name, variant1_name, metric, start_time, end_time)
    variant2_metrics = get_model_metrics(endpoint_name, variant2_name, metric, start_time, end_time)

    if metric == 'Invocation4XXErrors':
        # Lower errors are better
        better_variant = variant1_name if sum(variant1_metrics) < sum(variant2_metrics) else variant2_name
    else:
        # Higher throughput (Invocations) and lower latency (ModelLatency) are better
        better_variant = variant1_name if sum(variant1_metrics) > sum(variant2_metrics) else variant2_name

    print(f"{metric}: {better_variant} is performing better")

# Based on the overall comparison, make a decision on which variant is better

Invocation4XXErrors: Variant2-shadow is performing better
ModelLatency: Variant2-shadow is performing better
Invocations: Variant2-shadow is performing better

```

- **AS SEEN IN THE OUTPUT: shadow variant is better than production variant in terms of invocation and latency metrics.**
- **Shadow Test Completed with Success.**

Thank you